

Production-line workflow automation

Four-week cited memory pilot

Issue date: 2026-04-29

Prepared by: Daniel Panea Lichtig

Synthetic business document generated for the Private Company Memory Demo. All companies, people, domains, and scenarios are fictional and any resemblance to a real engagement is unintended.

Executive summary

Brannfeld Industrial GmbH is evaluating a focused private company memory layer for production-line workflow automation. The proposed pilot turns fragmented operational notes, CRM exports, and team documents into a cited account workspace that internal staff can query in plain language.

The strategic problem is recurring: relevant context exists somewhere in the company, but the people who need it before a customer-facing decision spend more time hunting for it than acting on it. Sales, operations, and customer-facing roles each rebuild the same context multiple times a week from email threads, attachments, and CRM exports, and the resulting briefings are inconsistent because each person stops at a different point.

Our recommendation is a bounded pilot with synthetic and approved business data first, followed by a measured expansion only after security review. The pilot is deliberately scoped to validate the architecture in a real environment without depending on production-grade connectors or autonomous actions, so the value can be evaluated against an honest cost.

The deliverable at the end of the pilot is a working demo on customer-controlled infrastructure, an evidence pack documenting retrieval quality and citation behaviour, and a rollout recommendation that names the next two or three concrete workflows that would justify a broader engagement.

Scope of work

The pilot configures ingestion for email archives, CRM exports, meeting notes, PDFs, and Word documents. Each artifact is parsed into raw extracted text, then normalised into AI-ready account memory documents with full source citations and retrieval metadata. The raw-artifact inventory remains visible behind every answer so internal reviewers can audit any claim.

Retrieval is hybrid: full-text search and dense embeddings are merged with reciprocal rank fusion, and an intent classifier picks the document types most likely to answer a given question. The classifier is conservative by design; when it cannot resolve which account a question refers to, the system asks rather than guessing.

Answer generation runs through a strict citation contract. Every substantive paragraph or bullet is paired with an exact source label that points back to the underlying artifact, and the system retries

when the model produces unsupported claims. Operators can inspect both the cited and the rejected attempts in the evidence pack.

Guided workflows are wired up for call briefing, recent-change review, open-risk surfacing, next-action recommendation, and follow-up drafting. These are entry points, not closed paths: every workflow opens a thread the user can extend with free-text follow-ups against the same retrieval context.

Capability	Pilot behavior
Source ingestion	Email archives, PDFs, Word documents, CRM CSV exports, and meeting notes.
Retrieval	Hybrid search over normalized account memory with visible source citations.
Answer generation	Guided workflows and free-text questions with citation validation.
Boundaries	No production automation, no autonomous actions, no broad SaaS feature set.

Risk and dependency notes

EU hosting and audit posture are explicit dependencies before production data is introduced.

Any production deployment would require customer-owned credentials, access review, logging policies, and a jointly approved retention policy. The pilot intentionally avoids touching production data so this approval can take its normal cadence without blocking the architecture review.

Two dependencies sit on the customer side and are worth flagging early. First, the source inventory needs at least one owner per artifact category — without that, the ingestion gets stuck on permission questions rather than on technical ones. Second, the evaluation framing must be agreed up front: a pilot that ends without a clear bar for success or failure tends to be relitigated rather than concluded.

On the architecture side, the main risk is over-broadening scope. The system intentionally does not orchestrate autonomous actions in the pilot; it produces cited briefings and drafts that a human approves. Where customer stakeholders ask for autonomous behaviour, that conversation is staged as a follow-on engagement with its own evaluation plan.

Timeline

Week 1 focuses on source inventory, access review, and success criteria. The goal is to leave the week with a shared list of artifact categories, named owners, and a written definition of what "good" looks like for the chosen workflows. No code change is more important than this alignment.

Week 2 builds the ingestion prototype and runs the first retrieval evaluation. Sample questions are drawn from the customer's actual workflows, and the evaluation captures retrieval quality, answer quality, and any citation failures, with the underlying retrieved context preserved for review.

Week 3 enables the guided workflows, hardens citation validation, and runs a stakeholder review with operations, IT, and any compliance contacts. The output is a working environment that internal reviewers can use, plus a written readout of the open issues and their mitigations.

Week 4 closes the pilot with a readout and a rollout recommendation. The recommendation explicitly names which next workflows would benefit, which integrations would have to be built, and which questions would still need to be answered before broader rollout. We treat "do not roll out" as an honest possible outcome.

Commercial summary

Item	Duration	Fee
Pilot implementation	4 weeks	EUR 85,000
Optional production readiness	Scoped later	Not included

Production-grade connector hardening, multi-tenant operations, billing, and admin dashboards are explicitly outside this pilot. They can be scoped as a separate engagement once the pilot has been reviewed.